

# Neural Networks - intro

## Part 1 - XOR

1. Using the XOR dataset below, train (400 epochs) a neural network (NN) using 1, 2, 3, 4, and 5 hidden layers (where each layer has only 2 neurons). For each n layers, store the resulting loss score along with n. Plot the results to find what the optimal number of layers is.
2. Repeat the above with 3 neurons in each Hidden layers. How do these results compare to the 2 neuron layers?
3. Repeat the above with 4 neurons in each Hidden layers. How do these results compare to the 2 and 3 neuron layers?
4. Using the most optimal configuraion (n-layers, k-neurons per layer), compare how `tanh`, `sigmoid`, `softplus` and `relu` effect the loss after 400 epochs. Try other Activation functions as well (<https://keras.io/activations/>)
5. Again with the most optimal setup, try other optimizers (instead of `SGD`) and report on the loss score. (<https://keras.io/optimizers/>)

## Part 2 - BYOD (Bring your own Dataset)

Using your own dataset, experiment and find the best Neural Network configuration. You may use any resource to improve results, just reference it.

While you may use any dataset, I'd prefer you didn't use the diabetes dataset used in the lesson.

<https://stackoverflow.com/questions/34673164/how-to-train-and-tune-an-artificial-multilayer-perceptron-neural-network-using-k>

<https://keras.io/>

```
In [1]: !pip3 install tensorflow keras
```

Defaulting to user installation because normal site-packages is not writeable  
Requirement already satisfied: tensorflow in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (2.8.0)  
Requirement already satisfied: keras in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (2.8.0)  
Requirement already satisfied: typing-extensions>=3.6.6 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (4.2.0)  
Requirement already satisfied: h5py>=2.9.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (3.6.0)  
Requirement already satisfied: tf-estimator-nightly==2.8.0.dev2021122109 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (2.8.0.dev2021122109)  
Requirement already satisfied: google-pasta>=0.1.1 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (0.2.0)  
Requirement already satisfied: termcolor>=1.1.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (1.1.0)  
Requirement already satisfied: six>=1.12.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (1.16.0)  
Requirement already satisfied: gast>=0.2.1 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (0.5.3)  
Requirement already satisfied: opt-einsum>=2.3.2 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (3.3.0)  
Requirement already satisfied: flatbuffers>=1.12 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (2.0)  
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (1.44.0)  
Requirement already satisfied: protobuf>=3.9.2 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (3.20.0)  
Requirement already satisfied: setuptools in /Library/Frameworks/Python.framework/Versions/3.10/lib/python3.10/site-packages (from tensorflow) (58.1.0)  
Requirement already satisfied: libclang>=9.0.1 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (13.0.0)  
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (0.24.0)  
Requirement already satisfied: tensorboard<2.9,>=2.8 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (2.8.0)  
Requirement already satisfied: absl-py>=0.4.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (1.0.0)  
Requirement already satisfied: astunparse>=1.6.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (1.6.3)  
Requirement already satisfied: numpy>=1.20 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (1.22.3)  
Requirement already satisfied: keras-preprocessing>=1.1.1 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (1.1.2)  
Requirement already satisfied: wrapt>=1.11.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (1.14.0)  
Requirement already satisfied: wheel<1.0,>=0.23.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from astunparse>=1.6.0->tensorflow) (0.37.1)  
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorboard<2.9,>=2.8->tensorflow) (0.6.1)  
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorboard<2.9,>=2.8->tensorflow) (0.4.6)  
Requirement already satisfied: requests<3,>=2.21.0 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorboard<2.9,>=2.8->tensorflow) (2.27.1)  
Requirement already satisfied: werkzeug>=0.11.15 in /Users/casemamamitsu/Library/Python/3.10/lib/python/site-packages (from tensorflow) (2.2.3)

ary/Python/3.10/lib/python/site-packages (from tensorboard<2.9,>=2.8->tensorflow) (2.1.1)

Requirement already satisfied: markdown>=2.6.8 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from tensorboard<2.9,>=2.8->tensorflow) (3.3.6)

Requirement already satisfied: google-auth<3,>=1.6.3 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from tensorboard<2.9,>=2.8->tensorflow) (2.6.5)

Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from tensorboard<2.9,>=2.8->tensorflow) (1.8.1)

Requirement already satisfied: pyasn1-modules>=0.2.1 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow) (0.2.8)

Requirement already satisfied: rsa<5,>=3.1.4 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow) (4.8)

Requirement already satisfied: cachetools<6.0,>=2.0.0 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow) (5.0.0)

Requirement already satisfied: requests-oauthlib>=0.7.0 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorboard<2.9,>=2.8->tensorflow) (1.3.1)

Requirement already satisfied: certifi>=2017.4.17 in /Library/Frameworks/Python.framework/Versions/3.10/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow) (2021.10.8)

Requirement already satisfied: idna<4,>=2.5 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow) (3.3)

Requirement already satisfied: charset-normalizer~=2.0.0 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow) (2.0.12)

Requirement already satisfied: urllib3<1.27,>=1.21.1 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow) (1.26.9)

Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow) (0.4.8)

Requirement already satisfied: oauthlib>=3.0.0 in /Users/casemymasamitsu/Library/Python/3.10/lib/python/site-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorboard<2.9,>=2.8->tensorflow) (3.2.0)

```
In [3]: from keras.models import Sequential
        from keras.layers import Dense
        from tensorflow.keras.optimizers import SGD #Stochastic Gradient Descent

        import numpy as np
        # fix random seed for reproducibility
        np.random.seed(7)


        import matplotlib.pyplot as plt
        %matplotlib inline
```

```
In [4]: n = 40
        xx = np.random.random((n,1))
        yy = np.random.random((n,1))
```

```
In [5]: X = np.array([np.array([xx,-xx,-xx,xx]),np.array([yy,-yy,yy,-yy])]).reshape(2,4)
        y = np.array([np.ones([2*n]),np.zeros([2*n])]).reshape(4*n)
```

In [6]: `plt.scatter(*zip(*X), c=y)`

Out[6]: `<matplotlib.collections.PathCollection at 0x29aa5a640>`



```
In [7]: num_layers = [1,2,3,4,5]
scores = []
for num_layer in num_layers:
    if num_layer == 1:
        model = Sequential()
        model.add(Dense(2, input_dim = 2, activation = 'tanh'))
        model.add(Dense(1, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=
        model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
    if num_layer == 2:
        model = Sequential()
        model.add(Dense(2, input_dim = 2, activation = 'tanh'))
        model.add(Dense(2, activation = 'tanh'))
        model.add(Dense(1, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=
        model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
    if num_layer == 3:
        model = Sequential()
        model.add(Dense(2, input_dim = 2, activation = 'tanh'))
        model.add(Dense(2, activation = 'tanh'))
        model.add(Dense(2, activation = 'tanh'))
        model.add(Dense(1, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=
        model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
    if num_layer == 4:
        model = Sequential()
        model.add(Dense(2, input_dim = 2, activation = 'tanh'))
        model.add(Dense(2, activation = 'tanh'))
        model.add(Dense(2, activation = 'tanh'))
        model.add(Dense(2, activation = 'tanh'))
        model.add(Dense(1, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=
        model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
    if num_layer == 5:
        model = Sequential()
```

```

model.add(Dense(2, input_dim = 2, activation = 'tanh'))
model.add(Dense(2, activation = 'tanh'))
model.add(Dense(2, activation = 'tanh'))
model.add(Dense(2, activation = 'tanh'))
model.add(Dense(2, activation = 'tanh'))
model.add(Dense(1, activation='sigmoid'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)

score = model.evaluate(X, y)
scores.append(score)

```

Metal device set to: Apple M1 Pro

```

2022-04-18 15:12:17.727905: I tensorflow/core/common_runtime/pluggable_device/
pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU
ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.
2022-04-18 15:12:17.728214: I tensorflow/core/common_runtime/pluggable_device/
pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/rep
lica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (devi
ce: 0, name: METAL, pci bus id: <undefined>)
2022-04-18 15:12:17.822187: W tensorflow/core/platform/profile_utils/cpu_util
s.cc:128] Failed to get CPU frequency: 0 Hz
2022-04-18 15:12:17.952926: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 4ms/step - loss: 0.4399 - accuracy:
0.6125

```

```

2022-04-18 15:13:33.681170: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-18 15:13:33.880658: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 3ms/step - loss: 0.3190 - accuracy:
0.8813

```

```

2022-04-18 15:14:59.645035: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-18 15:14:59.844993: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 4ms/step - loss: 0.3630 - accuracy:
0.6313

```

```

2022-04-18 15:16:30.153969: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-18 15:16:30.371859: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 4ms/step - loss: 0.3617 - accuracy:
0.7125

```

```

2022-04-18 15:18:07.273458: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-18 15:18:07.505004: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 4ms/step - loss: 0.3075 - accuracy:
0.8625

```

```

2022-04-18 15:19:51.706471: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

```


```

In [8]: columns = list(zip(*scores))
        loss = columns[0]

```

```
plt.scatter(num_layers, loss, c = "purple")
plt.title('Loss by # of Hidden Layers (1-5)')
plt.xlabel("# of Hidden Layers")
plt.ylabel('Loss')
```


Out[8]: Text(0, 0.5, 'Loss')



```
In [9]: accuracy = columns[1]

plt.scatter(num_layers, accuracy, c = "gold")
plt.title('Accuracy by # of Hidden Layers (1-5)')
plt.xlabel("# of Hidden Layers")
plt.ylabel('Accuracy')
```

Out[9]: Text(0, 0.5, 'Accuracy')



In both loss and accuracy, 3 hidden layers with 2 neurons is the optimal fit.

## Optimizers

Using the most optimal configuraion (n-layers, k-neurons per layer), compare how tanh, sigmoid,softplus and relu effect the loss after 400 epochs.

```
In [10]: # Tanh
model = Sequential()
model.add(Dense(2, input_dim = 2, activation = 'tanh'))
model.add(Dense(2, activation = 'tanh'))
model.add(Dense(2, activation = 'tanh'))
model.add(Dense(1, activation='tanh'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=['accuracy'])
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
tanh = model.evaluate(X, y)
```

```
2022-04-18 15:20:01.904793: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 5ms/step - loss: 3.8662 - accuracy: 0.7438
```

```
2022-04-18 15:21:33.816127: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
In [13]: # Sigmoid
model = Sequential()
model.add(Dense(2, input_dim = 2, activation = 'sigmoid'))
model.add(Dense(2, activation = 'sigmoid'))
model.add(Dense(2, activation = 'sigmoid'))
model.add(Dense(1, activation='sigmoid'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=['accuracy'])
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
sigmoid = model.evaluate(X, y)
```

```
2022-04-18 15:23:23.991514: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 4ms/step - loss: 0.6931 - accuracy: 0.5688
```

```
2022-04-18 15:24:52.088996: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
In [14]: # Softplus
model = Sequential()
model.add(Dense(2, input_dim = 2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(1, activation='softplus'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=['accuracy'])
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
softplus = model.evaluate(X, y)
```

```
2022-04-18 15:24:57.238863: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 7ms/step - loss: 0.2754 - accuracy: 0.8875
```

```
2022-04-18 15:28:00.595534: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
In [11]: # Relu
model = Sequential()
model.add(Dense(2, input_dim = 2, activation = 'relu'))
model.add(Dense(2, activation = 'relu'))
```



```
model.add(Dense(2, activation = 'relu'))
model.add(Dense(1, activation='relu'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=['accuracy'])
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
relu = model.evaluate(X, y)
```

```
2022-04-18 15:21:44.067677: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 3ms/step - loss: 7.7125 - accuracy: 0.5000
```

```
2022-04-18 15:23:14.186965: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
In [15]: print(tanh, sigmoid, softplus, relu)

[3.86621356010437, 0.7437500357627869] [0.6930978894233704, 0.5687500238418579] [0.2753864824771881, 0.887499988079071] [7.712474346160889, 0.5]
```

Of the four, softplus has the lowest loss and highest accuracy.

Try other Activation functions as well (<https://keras.io/activations/>) Again with the most optimal setup, try other optimizers (instead of SGD) and report on the loss score. (<https://keras.io/optimizers/>)

## Optimizing Softplus

```
In [16]: # Softplus with SGD:
model = Sequential()
model.add(Dense(2, input_dim = 2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(1, activation='softplus'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'sgd', metrics=['accuracy'])
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
sgd = model.evaluate(X, y)
```

```
2022-04-18 15:28:09.149143: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 7ms/step - loss: 0.2778 - accuracy: 0.8688
```

```
2022-04-18 15:31:08.734330: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
In [17]: # Softplus with RMSprop:
model = Sequential()
model.add(Dense(2, input_dim = 2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(1, activation='softplus'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'RMSprop', metrics=['accuracy'])
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
RMSprop = model.evaluate(X, y)
```



```
2022-04-18 15:31:12.638707: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 7ms/step - loss: 0.3095 - accuracy: 0.8813

2022-04-18 15:34:51.250138: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
In [18]: # Softplus with Adam:
model = Sequential()
model.add(Dense(2, input_dim = 2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(1, activation='softplus'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'adam', metrics=['accuracy'])
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
adam = model.evaluate(X, y)
```

```
2022-04-18 15:34:56.631691: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 7ms/step - loss: 0.3017 - accuracy: 0.8813

2022-04-18 15:38:16.705073: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
In [19]: # Softplus with Adagrad:
model = Sequential()
model.add(Dense(2, input_dim = 2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(2, activation = 'softplus'))
model.add(Dense(1, activation='softplus'))
sgd = SGD(learning_rate = 0.1)
model.compile(loss = 'binary_crossentropy', optimizer = 'adagrad', metrics=['accuracy'])
model.fit(X, y, batch_size = 2, epochs = 400, verbose = 0)
adagrad = model.evaluate(X, y)
```

```
2022-04-18 15:38:28.476550: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
5/5 [=====] - 0s 7ms/step - loss: 0.6904 - accuracy: 0.5563

2022-04-18 15:41:30.154040: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
In [20]: print(sgd, RMSprop, adam, adagrad)

<keras.optimizer_v2.gradient_descent.SGD object at 0x29f6fcbe0> [0.30945071578 0.8812500238418579] [0.30166906118392944, 0.8812500238418579] [0.690419 1970825195, 0.5562500357627869]
```

It is interesting, sometimes softplus "breaks" between 325-375 epochs and the loss score jumps from ~.50 to 7+. However, if I reduce the epochs to around 300, softplus is consistently the best. Seems there is a lot to play with here.

## Using Diabetes data

<http://archive.ics.uci.edu/ml/machine-learning-databases/pima-indians-diabetes/pima->

## indians-diabetes.data

1. Number of times pregnant
2. Plasma glucose concentration a 2 hours in an oral glucose tolerance test
3. Diastolic blood pressure (mm Hg)
4. Triceps skin fold thickness (mm)
5. 2-Hour serum insulin (mu U/ml)
6. Body mass index (weight in kg/(height in m)^2)
7. Diabetes pedigree function
8. Age (years)
9. Class variable (0 or 1)

```
In [21]: dataset = np.loadtxt("../data/pima-indians-diabetes.data", delimiter=",")
dataset
```

```
Out[21]: array([[ 6.   , 148.   , 72.   , ..., 0.627, 50.   , 1.   ],
 [ 1.   , 85.   , 66.   , ..., 0.351, 31.   , 0.   ],
 [ 8.   , 183.   , 64.   , ..., 0.672, 32.   , 1.   ],
 ...,
 [ 5.   , 121.   , 72.   , ..., 0.245, 30.   , 0.   ],
 [ 1.   , 126.   , 60.   , ..., 0.349, 47.   , 1.   ],
 [ 1.   , 93.   , 70.   , ..., 0.315, 23.   , 0.   ]])
```

```
In [23]: import pandas as pd
df = pd.DataFrame(dataset)
df
```

```
Out[23]:
```

	0	1	2	3	4	5	6	7	8
0	6.0	148.0	72.0	35.0	0.0	33.6	0.627	50.0	1.0
1	1.0	85.0	66.0	29.0	0.0	26.6	0.351	31.0	0.0
2	8.0	183.0	64.0	0.0	0.0	23.3	0.672	32.0	1.0
3	1.0	89.0	66.0	23.0	94.0	28.1	0.167	21.0	0.0
4	0.0	137.0	40.0	35.0	168.0	43.1	2.288	33.0	1.0
...	...	...	...	...	...	...	...	...	...
763	10.0	101.0	76.0	48.0	180.0	32.9	0.171	63.0	0.0
764	2.0	122.0	70.0	27.0	0.0	36.8	0.340	27.0	0.0
765	5.0	121.0	72.0	23.0	112.0	26.2	0.245	30.0	0.0
766	1.0	126.0	60.0	0.0	0.0	30.1	0.349	47.0	1.0
767	1.0	93.0	70.0	31.0	0.0	30.4	0.315	23.0	0.0

768 rows x 9 columns

```
In [24]: x = dataset[:,0:8]
print(x.shape)
x
```

(768, 8)

```
Out[24]: array([[ 6. , 148. , 72. , ..., 33.6 , 0.627, 50. ],
 [ 1. , 85. , 66. , ..., 26.6 , 0.351, 31. ],
 [ 8. , 183. , 64. , ..., 23.3 , 0.672, 32. ],
 ...,
 [ 5. , 121. , 72. , ..., 26.2 , 0.245, 30. ],
 [ 1. , 126. , 60. , ..., 30.1 , 0.349, 47. ],
 [ 1. , 93. , 70. , ..., 30.4 , 0.315, 23. ]])
```

```
In [25]: Y = dataset[:,8]
          print(Y.shape)
          Y
```

```
Out[25]: array([[1., 0., 1., 0., 1., 0., 1., 0., 1., 1., 0., 1., 0., 1., 1., 1., 1., 1.,
1., 0., 1., 0., 0., 1., 1., 1., 1., 1., 0., 0., 0., 0., 1., 0., 0.,
0., 0., 0., 1., 1., 1., 0., 0., 0., 1., 0., 0., 1., 0., 0., 0.,
0., 0., 1., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0., 1., 0., 1., 0.,
0., 0., 1., 0., 1., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 1.,
0., 0., 0., 1., 0., 0., 0., 0., 1., 1., 1., 0., 0., 1., 1., 0.,
0., 0., 0., 0., 0., 1., 1., 0., 0., 1., 1., 1., 1., 0., 0., 0.,
0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 1.,
0., 1., 1., 0., 0., 0., 1., 0., 0., 0., 0., 1., 1., 0., 0., 0., 0.,
1., 1., 0., 0., 0., 1., 0., 1., 0., 1., 0., 0., 0., 0., 1., 1.,
1., 1., 1., 0., 0., 1., 1., 0., 1., 0., 1., 1., 1., 0., 0., 0., 0.,
0., 0., 1., 1., 0., 1., 0., 0., 0., 1., 1., 1., 1., 0., 1., 1., 1.,
1., 0., 0., 0., 0., 0., 0., 1., 0., 0., 1., 1., 0., 0., 0., 1., 1., 1.,
1., 0., 0., 0., 1., 1., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 1.,
1., 0., 0., 0., 1., 0., 1., 0., 0., 1., 0., 0., 1., 1., 0., 0., 1.,
0., 0., 0., 0., 1., 0., 0., 0., 1., 0., 0., 1., 1., 0., 0., 1., 0.,
1., 1., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1.,
0., 1., 1., 1., 0., 0., 1., 0., 0., 1., 0., 0., 0., 1., 0., 0., 1.,
0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.,
1., 0., 0., 0., 0., 1., 0., 0., 0., 1., 1., 0., 0., 0., 0., 0., 0.,
1., 0., 0., 0., 0., 1., 1., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0.,
0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 1., 1., 1., 1., 0.,
0., 1., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1.,
1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1.,
1., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0.,
0., 1., 0., 0., 0., 0., 1., 0., 0., 1., 0., 0., 0., 0., 1., 0., 0.,
1., 1., 1., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 1., 0., 1., 1.,
0., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 1., 1., 0., 0., 1., 0.,
0., 1., 1., 0., 0., 1., 0., 0., 1., 0., 0., 0., 0., 0., 0., 1.,
1., 1., 0., 0., 0., 0., 0., 0., 1., 1., 0., 0., 1., 0., 0., 1., 0.,
1., 1., 1., 0., 0., 1., 1., 1., 0., 1., 0., 1., 0., 0., 0., 0.,
0., 1., 0.]])
```

```
In [26]: # create model
model = Sequential()
model.add(Dense(16, input_dim=8, activation='tanh'))
model.add(Dense(16, activation='tanh'))
model.add(Dense(1, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, Y, epochs=1000, batch_size=10, verbose = 0)
# evaluate the model
scores = model.evaluate(X, Y)
print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

```
2022-04-18 15:41:55.459088: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
24/24 [=====] - 0s 3ms/step - loss: 0.4384 - accuracy: 0.7917
```

accuracy: 79.17%

```
2022-04-18 15:47:15.337316: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

## Part 2: BYOD - Wine Dataset

```
In [27]: import pandas as pd
wine = pd.read_csv('../data/WineQT.csv', names=['fixedacid', 'volatileacid', 'citricacid', 'residualsugar', 'chlorides', 'freesulfurdio', 'totalsulfurdio'])
wine = wine.iloc[1: , :]
wine
```

```
Out[27]:
```

	fixedacid	volatileacid	citricacid	residualsugar	chlorides	freesulfurdio	totalsulfurdio
1	7.4	0.7	0	1.9	0.076	11	34
2	7.8	0.88	0	2.6	0.098	25	67
3	7.8	0.76	0.04	2.3	0.092	15	54
4	11.2	0.28	0.56	1.9	0.075	17	60
5	7.4	0.7	0	1.9	0.076	11	34
...	...	...	...	...	...	...	...
1138	5.4	0.74	0.09	1.7	0.089	16	26
1139	6.3	0.51	0.13	2.3	0.076	29	40
1140	6.8	0.62	0.08	1.9	0.068	28	38
1141	6.2	0.6	0.08	2	0.09	32	44
1142	5.9	0.55	0.1	2.2	0.062	39	51

1142 rows x 12 columns

```
In [28]: wine["fixedacid"] = pd.to_numeric(wine.fixedacid, errors='coerce')
wine["volatileacid"] = pd.to_numeric(wine.volatileacid, errors='coerce')
wine["citricacid"] = pd.to_numeric(wine.citricacid, errors='coerce')
wine["residualsugar"] = pd.to_numeric(wine.residualsugar, errors='coerce')
wine["chlorides"] = pd.to_numeric(wine.chlorides, errors='coerce')
```

```
wine["freesulfurdio"] = pd.to_numeric(wine.freesulfurdio, errors='coerce')
wine["totalsulfurdio"] = pd.to_numeric(wine.totalsulfurdio, errors='coerce')
wine["density"] = pd.to_numeric(wine.density, errors='coerce')
wine["pH"] = pd.to_numeric(wine.pH, errors='coerce')
wine["sulphates"] = pd.to_numeric(wine.sulphates, errors='coerce')
wine["alcohol"] = pd.to_numeric(wine.alcohol, errors='coerce')
wine["quality"] = pd.to_numeric(wine.quality, errors='coerce')
```

```
In [29]: wine.head()
print(wine.dtypes)
```

```
fixedacid          float64
volatileacid       float64
citricacid         float64
residualsugar      float64
chlorides          float64
freesulfurdio      float64
totalsulfurdio     float64
density           float64
pH               float64
sulphates         float64
alcohol          float64
quality          int64
dtype: object
```

```
In [30]: wine_array = wine.to_numpy()
wine_array
```

```
Out[30]: array([[ 7.4 ,  0.7 ,  0.  , ...,  0.56,  9.4 ,  5.  ],
 [ 7.8 ,  0.88,  0.  , ...,  0.68,  9.8 ,  5.  ],
 [ 7.8 ,  0.76,  0.04, ...,  0.65,  9.8 ,  5.  ],
 ...,
 [ 6.8 ,  0.62,  0.08, ...,  0.82,  9.5 ,  6.  ],
 [ 6.2 ,  0.6 ,  0.08, ...,  0.58, 10.5 ,  5.  ],
 [ 5.9 ,  0.55,  0.1 , ...,  0.76, 11.2 ,  6.  ]])
```

```
In [31]: X = wine_array[:,0:12]
print(X.shape)
X
```

```
Out[31]: (1142, 12)
array([[ 7.4 ,  0.7 ,  0.  , ...,  0.56,  9.4 ,  5.  ],
 [ 7.8 ,  0.88,  0.  , ...,  0.68,  9.8 ,  5.  ],
 [ 7.8 ,  0.76,  0.04, ...,  0.65,  9.8 ,  5.  ],
 ...,
 [ 6.8 ,  0.62,  0.08, ...,  0.82,  9.5 ,  6.  ],
 [ 6.2 ,  0.6 ,  0.08, ...,  0.58, 10.5 ,  5.  ],
 [ 5.9 ,  0.55,  0.1 , ...,  0.76, 11.2 ,  6.  ]])
```

```
In [32]: quality = pd.get_dummies(wine["quality"])
```

```
In [33]: Y = pd.get_dummies(wine["quality"])
print(Y.shape)
Y
```

```
(1142, 6)
```

Out[33]:

	3	4	5	6	7	8
1	0	0	1	0	0	0
2	0	0	1	0	0	0
3	0	0	1	0	0	0
4	0	0	0	1	0	0
5	0	0	1	0	0	0
...	...	...	...	...	...	...
1138	0	0	0	1	0	0
1139	0	0	0	1	0	0
1140	0	0	0	1	0	0
1141	0	0	1	0	0	0
1142	0	0	0	1	0	0

1142 rows × 6 columns

```
In [34]: # Create model
model = Sequential()
model.add(Dense(20, input_dim = 12, activation='tanh'))
model.add(Dense(20, input_dim = 12, activation='tanh'))
model.add(Dense(20, input_dim = 12, activation='tanh'))
model.add(Dense(6, activation='sigmoid'))

# Compile model
sgd = SGD(learning_rate = 0.1)
model.compile(loss='CategoricalCrossentropy', optimizer='sgd', metrics=['accuracy'])

# Fit the model
model.fit(X, Y, epochs = 500, batch_size = 10, verbose = 0)

# Evaluate the model
scores = model.evaluate(X, Y)
print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

```
2022-04-18 15:47:26.330155: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
29/36 [=====>.....] - ETA: 0s - loss: 0.5612 - accuracy: 0.7823
```

```
2022-04-18 15:52:58.437220: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
36/36 [=====] - 0s 6ms/step - loss: 0.5576 - accuracy: 0.7890
```

accuracy: 78.90%

After a while playing around with activations, optimizers, and neurons above, I found the best results with 20 neurons, tanh/sigmoid, and sgd as an optimizer. Below, I ran the same code as above to test it out with 1-5 layers.

```
In [35]: num_layers = [1,2,3,4,5]
```

```

scores = []
for num_layer in num_layers:
    if num_layer == 1:
        model = Sequential()
        model.add(Dense(20, input_dim = 12, activation = 'tanh'))
        model.add(Dense(6, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'CategoricalCrossentropy', optimizer = 'sgd', metrics=['accuracy'])
        model.fit(X, Y, batch_size = 10, epochs = 300, verbose = 0)
    if num_layer == 2:
        model = Sequential()
        model.add(Dense(20, input_dim = 12, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(6, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'CategoricalCrossentropy', optimizer = 'sgd', metrics=['accuracy'])
        model.fit(X, Y, batch_size = 10, epochs = 300, verbose = 0)
    if num_layer == 3:
        model = Sequential()
        model.add(Dense(20, input_dim = 12, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(6, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'CategoricalCrossentropy', optimizer = 'sgd', metrics=['accuracy'])
        model.fit(X, Y, batch_size = 10, epochs = 300, verbose = 0)
    if num_layer == 4:
        model = Sequential()
        model.add(Dense(20, input_dim = 12, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(6, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'CategoricalCrossentropy', optimizer = 'sgd', metrics=['accuracy'])
        model.fit(X, Y, batch_size = 10, epochs = 300, verbose = 0)
    if num_layer == 5:
        model = Sequential()
        model.add(Dense(20, input_dim = 12, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(20, activation = 'tanh'))
        model.add(Dense(6, activation='sigmoid'))
        sgd = SGD(learning_rate = 0.1)
        model.compile(loss = 'CategoricalCrossentropy', optimizer = 'sgd', metrics=['accuracy'])
        model.fit(X, Y, batch_size = 10, epochs = 300, verbose = 0)

score = model.evaluate(X, Y)
scores.append(score)

```

```

2022-04-18 15:53:48.443395: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
30/36 [=====>.....] - ETA: 0s - loss: 0.6000 - accuracy: 0.8063

```

```

2022-04-18 15:56:47.320419: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

```



36/36 [=====] - 0s 6ms/step - loss: 0.5990 - accuracy: 0.8135

2022-04-18 15:56:47.681272: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

29/36 [=====>.....] - ETA: 0s - loss: 0.8295 - accuracy: 0.6584

2022-04-18 15:59:53.576048: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

36/36 [=====] - 0s 6ms/step - loss: 0.8212 - accuracy: 0.6602

2022-04-18 15:59:53.963050: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

30/36 [=====>.....] - ETA: 0s - loss: 0.4454 - accuracy: 0.8635

2022-04-18 16:03:11.149469: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

36/36 [=====] - 0s 6ms/step - loss: 0.4272 - accuracy: 0.8687

2022-04-18 16:03:11.544112: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

29/36 [=====>.....] - ETA: 0s - loss: 0.3286 - accuracy: 0.9106

2022-04-18 16:06:45.120456: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

36/36 [=====] - 0s 6ms/step - loss: 0.3287 - accuracy: 0.9116

2022-04-18 16:06:45.540390: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

26/36 [=====>.....] - ETA: 0s - loss: 0.3522 - accuracy: 0.8702

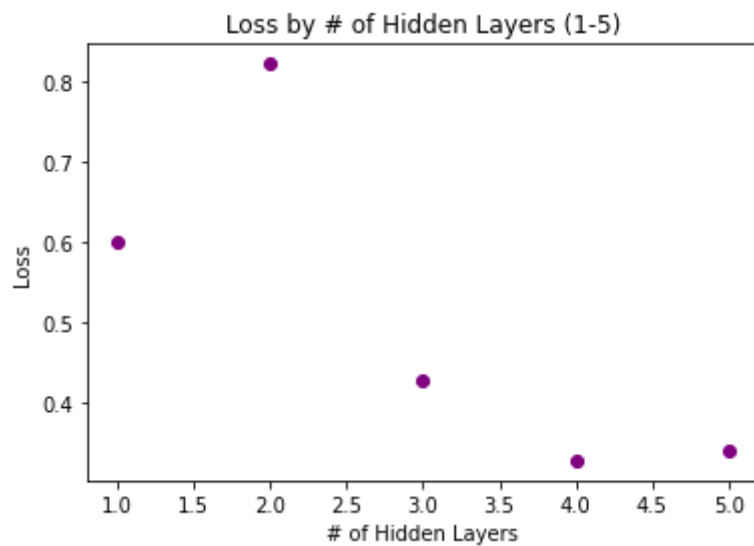
2022-04-18 16:10:39.376238: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

36/36 [=====] - 0s 6ms/step - loss: 0.3406 - accuracy: 0.8783

```
In [36]: columns = list(zip(*scores))
         loss = columns[0]

         plt.scatter(num_layers, loss, c = "purple")
         plt.title('Loss by # of Hidden Layers (1-5)')
         plt.xlabel("# of Hidden Layers")
         plt.ylabel('Loss')
```


```
Out[36]: Text(0, 0.5, 'Loss')
```



```
In [37]: accuracy = columns[1]

plt.scatter(num_layers, accuracy, c = "gold")
plt.title('Accuracy by # of Hidden Layers (1-5)')
plt.xlabel("# of Hidden Layers")
plt.ylabel('Accuracy')
```

```
Out[37]: Text(0, 0.5, 'Accuracy')
```



Four hidden layers with 20 neurons each had the highest accuracy and lowest loss.

```
In [ ]:
```