

# Assignment is below at the bottom

Video 13.1 <https://www.youtube.com/watch?v=kIGHE7Cfe1s>

Video 13.2 <https://www.youtube.com/watch?v=Rm9bJcDd1KU>

Video 13.3 <https://youtu.be/6HjZk-3LsjE>

```
In [1]: import numpy as np
import tensorflow as tf
```

```
In [10]: from keras.callbacks import TensorBoard
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import mnist
import numpy as np
import matplotlib.pyplot as plt

(xtrain, ytrain), (xtest, ytest) = mnist.load_data()

xtrain = xtrain.astype('float32') / 255.
xtest = xtest.astype('float32') / 255.
xtrain = xtrain.reshape((len(xtrain), np.prod(xtrain.shape[1:]))
xtest = xtest.reshape((len(xtest), np.prod(xtest.shape[1:]))
xtrain.shape, xtest.shape
```

```
Out[10]: ((60000, 784), (10000, 784))
```

```
In [11]: # this is the size of our encoded representations
encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input

# this is our input placeholder
x = input_img = Input(shape=(784,))

# "encoded" is the encoded representation of the input
x = Dense(256, activation='relu')(x)
x = Dense(128, activation='relu')(x)
encoded = Dense(encoding_dim, activation='relu')(x)

# "decoded" is the lossy reconstruction of the input
x = Dense(128, activation='relu')(encoded)
x = Dense(256, activation='relu')(x)
decoded = Dense(784, activation='sigmoid')(x)

# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)

encoder = Model(input_img, encoded)

# create a placeholder for an encoded (32-dimensional) input
encoded_input = Input(shape=(encoding_dim,))
# retrieve the last layer of the autoencoder model
dcd1 = autoencoder.layers[-1]
dcd2 = autoencoder.layers[-2]
dcd3 = autoencoder.layers[-3]
```

```
# create the decoder model
decoder = Model(encoded_input, dcd1(dcd2(dcd3(encoded_input))))
```

```
In [12]: autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
```

```
In [13]: autoencoder.fit(xtrain, xtrain,
                        epochs=10,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(xtest, xtest))
```

Epoch 1/10

9/235 [>.....] - ETA: 1s - loss: 0.6935

2022-04-20 19:05:16.565500: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

235/235 [=====] - 2s 7ms/step - loss: 0.6934 - val\_loss: 0.6934

Epoch 2/10

1/235 [.....] - ETA: 1s - loss: 0.6934

2022-04-20 19:05:18.093062: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

235/235 [=====] - 1s 6ms/step - loss: 0.6933 - val\_loss: 0.6933

Epoch 3/10

235/235 [=====] - 1s 6ms/step - loss: 0.6932 - val\_loss: 0.6932

Epoch 4/10

235/235 [=====] - 1s 6ms/step - loss: 0.6931 - val\_loss: 0.6931

Epoch 5/10

235/235 [=====] - 2s 6ms/step - loss: 0.6930 - val\_loss: 0.6930

Epoch 6/10

235/235 [=====] - 1s 6ms/step - loss: 0.6929 - val\_loss: 0.6929

Epoch 7/10

235/235 [=====] - 1s 6ms/step - loss: 0.6928 - val\_loss: 0.6928

Epoch 8/10

235/235 [=====] - 1s 6ms/step - loss: 0.6927 - val\_loss: 0.6927

Epoch 9/10

235/235 [=====] - 1s 6ms/step - loss: 0.6926 - val\_loss: 0.6926

Epoch 10/10

235/235 [=====] - 1s 6ms/step - loss: 0.6925 - val\_loss: 0.6925

```
Out[13]: <keras.callbacks.History at 0x2b8d6afa0>
```

```
In [7]: encoded_imgs = encoder.predict(xtest)
        decoded_imgs = decoder.predict(encoded_imgs)
        import matplotlib.pyplot as plt

        n = 20 # how many digits we will display
        plt.figure(figsize=(40, 4))
        for i in range(n):
            # display original
```

```

ax = plt.subplot(2, n, i + 1)
plt.imshow(xtest[i].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)


# display reconstruction
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(decoded_imgs[i].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()

```

```

2022-04-20 19:03:50.025844: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:03:50.377605: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

```



```

In [15]: noise = np.random.normal(20,4, (4,32))
noise_preds = decoder.predict(noise)

```

```

2022-04-20 19:05:58.056183: I tensorflow/core/grappler/optimizers/custom_graph
_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

```

```


In [16]: plt.imshow(noise_preds[1].reshape(28,28))

```

```

Out[16]: <matplotlib.image.AxesImage at 0x2c44b9d00>

```



```

In [17]: np.max(encoded_imgs)

```

```

Out[17]: 1.3580321

```

```

In [18]: encoded_imgs


```

```
Out[18]: array([[0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ..., 2.3225811e-01,
                0.00000000e+00, 2.5551948e-01],
                [0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ..., 1.7010622e-01,
                4.1145870e-01, 9.5288701e-02],
                [0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ..., 4.2720996e-02,
                0.00000000e+00, 3.2465896e-01],
                ...,
                [0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ..., 0.00000000e+00,
                4.1268975e-02, 3.0748990e-01],
                [0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ..., 9.7543161e-05,
                9.6587375e-02, 2.7234137e-01],
                [0.00000000e+00, 0.00000000e+00, 0.00000000e+00, ..., 3.0143291e-01,
                3.2321402e-01, 5.5349982e-01]], dtype=float32)
```

```
In [19]: %matplotlib inline
```


```
In [20]: plt.scatter(encoded_imgs[:,1], encoded_imgs[:,0], s=1, c=ytest, cmap='rainbow')
        # plt.show()
```

```
Out[20]: <matplotlib.collections.PathCollection at 0x2b8dfd2e0>
```



```
In [21]: plt.scatter(encoded_imgs[:,1], encoded_imgs[:,3], s=1, c=ytest, cmap='rainbow')
        # plt.show()
```


```
Out[21]: <matplotlib.collections.PathCollection at 0x2cffb7f40>
```



```
In [22]: plt.scatter(encoded_imgs[:,1], encoded_imgs[:,2], s=1, c=ytest, cmap='rainbow')
```


```
# plt.show()
```

Out[22]: <matplotlib.collections.PathCollection at 0x2c4456070>



```
In [23]: from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(encoded_imgs[:,1], encoded_imgs[:,2], encoded_imgs[:,3], c=ytest, cn
```

Out[23]: <mpl\_toolkits.mplot3d.art3d.Path3DCollection at 0x17f7a8130>



In [ ]:

## Assignment

1. change the `encoding_dim` through various values ( `range(2,18,2)` ) and store or keep track of the best loss you can get. Plot the 8 pairs of dimensions vs loss on a scatter plot

```
In [26]: losses = []
encoding_dim = (range(2, 18, 2))

encoded_imgs = encoder.predict(xtest)
decoded_imgs = decoder.predict(encoded_imgs)
```

```

for encoding_dim in encoding_dim:

    x = input_img = Input(shape=(784,))

    x = Dense(256, activation='relu')(x)
    x = Dense(128, activation='relu')(x)
    encoded = Dense(encoding_dim, activation='relu')(x)

    x = Dense(128, activation='relu')(encoded)
    x = Dense(256, activation='relu')(x)
    decoded = Dense(784, activation='sigmoid')(x)

    autoencoder = Model(input_img, decoded)
    encoder = Model(input_img, encoded)

    encoded_input = Input(shape=(encoding_dim,))

    dcd1 = autoencoder.layers[-1]
    dcd2 = autoencoder.layers[-2]
    dcd3 = autoencoder.layers[-3]

    decoder = Model(encoded_input, dcd1(dcd2(dcd3(encoded_input))))

    autoencoder.compile(optimizer=tf.keras.optimizers.Adadelta(learning_rate=1))

    model = autoencoder.fit(xtrain, xtrain,
                            epochs=100,
                            batch_size=256,
                            shuffle=True,
                            verbose=0,
                            validation_data=(xtest, xtest))

    n = 20
    plt.figure(figsize=(60, 4))
    for i in range(n):
        ax = plt.subplot(2, n, i + 1)
        plt.imshow(xtest[i].reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

        ax = plt.subplot(2, n, i + 1 + n)
        plt.imshow(decoded_imgs[i].reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)
    plt.show()

    loss = np.min(model.history["loss"])
    losses.append(loss)

```

```
2022-04-20 19:10:32.098635: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:10:32.449981: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:10:33.144721: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:10:34.562938: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

7	2	1	0	4	1	4	9	5	9	0	6	9	0	1	5	9	7	8	4
7	8	1	0	9	1	9	9	8	7	0	6	9	0	1	8	9	7	6	9

```
2022-04-20 19:12:42.101737: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:12:43.612989: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

7	2	1	0	4	1	4	9	5	9	0	6	9	0	1	5	9	7	8	4
7	8	1	0	9	1	9	9	8	7	0	6	9	0	1	8	9	7	6	9

```
2022-04-20 19:14:52.686949: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:14:54.034494: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

7	2	1	0	4	1	4	9	5	9	0	6	9	0	1	5	9	7	8	4
7	8	1	0	9	1	9	9	8	7	0	6	9	0	1	8	9	7	6	9

```
2022-04-20 19:17:04.097148: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:17:05.851535: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

7	2	1	0	4	1	4	9	5	9	0	6	9	0	1	5	9	7	8	4
7	8	1	0	9	1	9	9	8	7	0	6	9	0	1	8	9	7	6	9

```
2022-04-20 19:19:16.498731: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:19:17.919002: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

7	2	1	0	4	1	4	9	5	9	0	6	9	0	1	5	9	7	8	4
7	8	1	0	9	1	9	9	8	7	0	6	9	0	1	8	9	7	6	9

```
2022-04-20 19:21:26.398818: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:21:27.876831: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

7	2	1	0	4	1	4	9	5	9	0	6	9	0	1	5	9	7	8	4
7	8	1	0	9	1	9	9	8	7	0	6	9	0	1	8	9	7	6	9

```
2022-04-20 19:23:33.971330: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:23:35.378756: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

7	2	1	0	4	1	4	9	5	9	0	6	9	0	1	5	9	7	8	4
7	8	1	0	9	1	9	9	8	7	0	6	9	0	1	8	9	7	6	9

```
2022-04-20 19:25:41.635916: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:25:43.075658: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```



```
In [27]: encoding_dim = np.asarray(range(2, 18, 2))
         losses = np.asarray(losses)
```

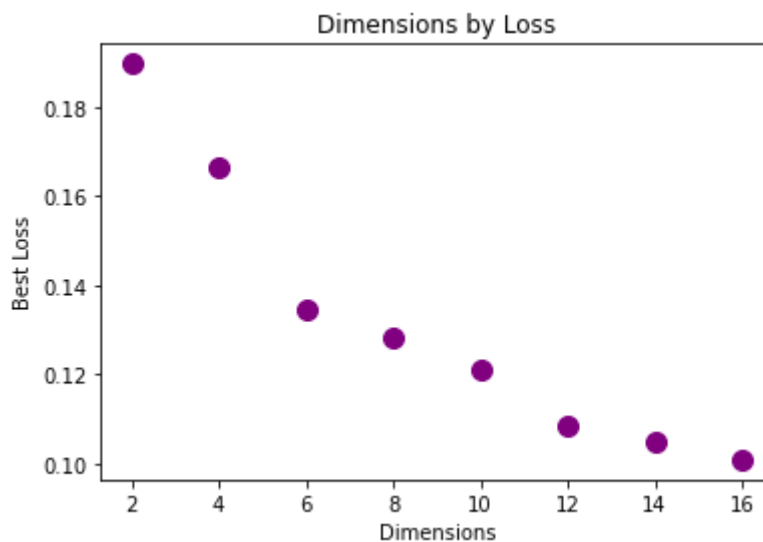
```
In [28]: encoding_dim.shape, losses.shape
```

```
Out[28]: ((8,), (8,))
```

```
In [29]: encoded_imgs
```

```
Out[29]: array([[ 4.2861347, 22.637783,  0.        , 18.613836 ],
                [ 9.403811,  7.13039,  0.        ,  6.5799375],
                [24.402075, 40.785385,  0.        , 13.966941 ],
                ...,
                [ 4.916314, 16.677427,  0.        ,  8.452135 ],
                [ 6.9526405,  7.2563763,  0.        ,  0.82639 ],
                [20.648458,  9.083031,  0.        , 11.402394 ]], dtype=float32)
```

```
In [30]: plt.scatter(encoding_dim, losses, s=100, c = "purple")
         plt.title("Dimensions by Loss")
         plt.xlabel("Dimensions")
         plt.ylabel("Best Loss")
         plt.show()
```



1. **After** training an autoencoder with `encoding_dim=8`, apply noise (like the previous assignment) to *only* the input of the trained autoencoder (not the output). The output images should be without noise.

```
In [31]: (xtrain, ytrain), (xtest, ytest) = mnist.load_data()

         xtrain = xtrain.astype('float32') / 255.
         xtest = xtest.astype('float32') / 255.
         xtrain = xtrain.reshape((len(xtrain), np.prod(xtrain.shape[1:])))
         xtest = xtest.reshape((len(xtest), np.prod(xtest.shape[1:])))
         xtrain.shape, xtest.shape
```



Out[31]: ((60000, 784), (10000, 784))

```
In [32]: encoding_dim = 8

x = input_img = Input(shape=(784,))

x = Dense(256, activation='relu')(x)
x = Dense(128, activation='relu')(x)
encoded = Dense(encoding_dim, activation='relu')(x)

x = Dense(128, activation='relu')(encoded)
x = Dense(256, activation='relu')(x)
decoded = Dense(784, activation='sigmoid')(x)

autoencoder = Model(input_img, decoded)

encoder = Model(input_img, encoded)

encoded_input = Input(shape=(encoding_dim,))

dcd1 = autoencoder.layers[-1]
dcd2 = autoencoder.layers[-2]
dcd3 = autoencoder.layers[-3]

decoder = Model(encoded_input, dcd1(dcd2(dcd3(encoded_input))))
```

```
In [34]: autoencoder.compile(optimizer=tf.keras.optimizers.Adadelta(learning_rate=1), loss='binary_crossentropy')
autoencoder.fit(xtrain, xtrain,
                epochs=100,
                batch_size=256,
                shuffle=True,
                verbose=0,
                validation_data=(xtest, xtest))
```

```
2022-04-20 19:30:00.776915: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:30:02.122866: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

Out[34]: <keras.callbacks.History at 0x2eb1c69d0>

```
In [35]: encoded_imgs = encoder.predict(xtest)
decoded_imgs = decoder.predict(encoded_imgs)
import matplotlib.pyplot as plt

n = 20 # how many digits we will display
plt.figure(figsize=(40, 4))
for i in range(n):
    # display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(xtest[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
```

```
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()
```

```
2022-04-20 19:33:05.093858: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2022-04-20 19:33:05.521201: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```



```
In [80]: print(np.min(encoded_imgs))
print(np.max(encoded_imgs))
```

```
0.0
42.548054
```

```
In [134... noise = np.random.normal(20,5, (10000,784))
noise_preds = autoencoder.predict(noise)
```


```
In [135... noise_preds.shape
```

```
Out[135]: (10000, 784)
```

Print a few noisy images along with the output images to show they don't have noise.

```
In [136... plt.imshow(noise_preds[0].reshape(28,28))
```

```
Out[136]: <matplotlib.image.AxesImage at 0x162164880>
```




```
In [137... decoded_imgs.shape
```

```
Out[137]: (10000, 784)
```


```
In [138... plt.imshow(decoded_imgs[0].reshape(28,28))
```

```
Out[138]: <matplotlib.image.AxesImage at 0x161f26730>
```




```
In [55]: plt.imshow(noise_preds[1].reshape(28, 28))
```

```
Out[55]: <matplotlib.image.AxesImage at 0x167405670>
```




```
In [131]: plt.imshow(decoded_imgs[1].reshape(28, 28))
```

```
Out[131]: <matplotlib.image.AxesImage at 0x161eab850>
```




```
In [59]: plt.imshow(noise_preds[2].reshape(28, 28))
```

```
Out[59]: <matplotlib.image.AxesImage at 0x2cb277f70>
```



```
In [132... plt.imshow(decoded_imgs[2].reshape(28,28))
```

Out[132]: <matplotlib.image.AxesImage at 0x161ed7e20>



In [ ]: